

Thesis Simulation Single Run for Chapter 4

Dasha Asienaga

2024-03-25

Contents

Logistic Regression	1
Reading in the Data	1
Fitting the Logistic Regression	1
Convergence	2
Accuracy	2
Discrimination	2
Seldonian Framework	3
Reading in the Data	3
Pre-Processing the Data	3
Fitting a Seldonian Algorithm	4
Convergence	7
Accuracy	8
Discrimination	8

This file is intended to run the simulation process on the parent simulation data set as a single run, before scaling it into multiple trials.

Logistic Regression

Reading in the Data

First, let's read in the parent simulation data set.

```
compas_sim_path <- "/home/dasienga24/Statistics-Senior-Honors-Thesis/Data Sets/COMPAS/compas_sim.csv"
compas_sim_parent <- read.csv(compas_sim_path)
```

Fitting the Logistic Regression

We'll want to run the logistic regression as our baseline model and obtain the 3 key performance measures: convergence, accuracy, and discrimination.

```
lr <- glm(is_recid ~ age + prior_offense,
          data = compas_sim_parent,
          family = binomial(logit))
```

Convergence

Obtain convergence as an object. We would expect LR to always converge.

```
lr_converged <- lr[["converged"]]  
lr_converged
```

```
## [1] TRUE
```

Accuracy

Conditional on convergence, obtain the accuracy as an object, which in this case is 60.76%.

```
lr_accuracy <- count(round(lr[["fitted.values"]]) == lr[["y"]])/nrow(compas_sim_parent)  
lr_accuracy
```

```
## n_TRUE  
## 0.6076
```

Discrimination

Conditional on convergence, obtain the discrimination statistic as an object, which in this case is 0.2784 or 27.84%.

```
preds <- predict(lr, newdata = compas_sim_parent, type="response")  
  
compas_sim_parent <- compas_sim_parent %>%  
  mutate(preds = preds,  
         prediction = round(preds, 0),  
         pred_risk = ifelse(prediction == 0, 'Low', 'High'))  
  
discrimination <- compas_sim_parent %>%  
  dplyr::select(race, pred_risk, is_recid) %>%  
  group_by(race, is_recid) %>%  
  mutate(total = n()) %>%  
  group_by(pred_risk, race, total) %>%  
  summarise("reoffended" = count(is_recid == 1),  
           "did_not_reoffend" = count(is_recid == 0)) %>%  
  pivot_longer(cols = c("reoffended", "did_not_reoffend"),  
              names_to = "recidivism") %>%  
  pivot_wider(  
    id_cols = c("pred_risk", "recidivism", "total"),  
    names_from = "race",  
    values_from = value  
  ) %>%  
  rename("Black" = `African-American`,  
        "White" = `Caucasian`) %>%  
  mutate(Black = round(100 * Black / total, 2),  
         White = round(100 * White / total, 2)) %>%  
  dplyr::select(-total) %>%  
  group_by(pred_risk, recidivism) %>%  
  summarize(Black = max(Black, na.rm = TRUE),
```

```

        White = max(White, na.rm = TRUE)) %>%
filter((pred_risk == "High" & recidivism == "did_not_reoffend") |
       (pred_risk == "Low" & recidivism == "reoffended"))
)

lr_disc_stat <- sum(abs(discrimination$White - discrimination$Black))/100
lr_disc_stat

## [1] 0.2784

```

The results from the logistic regression are now easily savable and retrievable, which will be useful when we scale the process.

Seldonian Framework

We also want to be able to easily retrieve the 3 key performance measures from the Seldonian algorithms we run before we scale the process, that is, convergence, accuracy, and discrimination.

Reading in the Data

First, let's read in the data.

```

# point to the data file
f_orig = "/home/dasienga24/Statistics-Senior-Honors-Thesis/Data Sets/COMPAS/compas_sim.csv"

columns_orig = ["race", "prior_offense", "age", "is_recid"]

df = pd.read_csv(f_orig, header=0, names=columns_orig)

```

Pre-Processing the Data

Let's also preprocess the data to prepare it for Seldonian modeling.

```

# select inputs to be transformed
X = df.drop(columns=["is_recid"])
y = df["is_recid"]

# one hot encode race, scale age using standard scaler
ct = ColumnTransformer([('c', OneHotEncoder(), ['race']), ('n', StandardScaler(), ['age'])])

# apply transformation
X_transformed = ct.fit_transform(X)

# get names after one-hot encoding
output_columns = ct.get_feature_names_out(ct.feature_names_in_)

# make an output dataframe to save transformed X and y
outdf = pd.DataFrame(X_transformed, columns=output_columns)

# change names of columns
outdf.rename(columns={'c__race_African-American': 'Black', 'c__race_Caucasian': 'White', 'n__age': 'age'},

```

```

# re-index in order to concatenate columns
prior_offense = df["prior_offense"]
y.index = range(0, len(y))
prior_offense.index = range(0, len(prior_offense))

# add label column and `prior_offense` into final dataframe
outdf['prior_offense'] = prior_offense
outdf['is_recid'] = y

```

The data set is now clean and ready. Let's save it along with the JSON metadata file.

```

# save final dataframe
output_path_data="/home/dasienga24/Statistics-Senior-Honors-Thesis/Data Sets/COMPAS/Simulation Single Run/"

outdf.to_csv(output_path_data,index=False,header=False)

# save metadata json file
output_path_metadata="/home/dasienga24/Statistics-Senior-Honors-Thesis/Data Sets/COMPAS/Simulation Single Run/"

metadata_dict = {
    "regime":"supervised_learning",
    "sub_regime":"classification",
    "all_col_names":list(outdf.columns),
    "label_col_names":"is_recid",
    "sensitive_col_names":["Black", "White"]
}

with open(output_path_metadata,'w') as outfile:
    json.dump(metadata_dict,outfile,indent=2)

```

Fitting a Seldonian Algorithm

Varying $\epsilon = 0.2, 0.1, 0.05$, & 0.01 , let's fit a Seldonian algorithm such that

$$abs((FNR|[Black]) - (FNR|[White])) + abs((FPR|[Black]) - (FPR|[White])) \leq \epsilon.$$

We will take $\delta = 0.05$ to ensure 95% confidence.

First, let's read in the data set and specify the regime.

```

import autograd.numpy as np

data_pth = output_path_data
metadata_pth = output_path_metadata
save_dir = "/home/dasienga24/Statistics-Senior-Honors-Thesis/Python/COMPAS Simulation/Single Run/"
os.makedirs(save_dir,exist_ok=True)

# create dataset from data and metadata file
regime='supervised_learning'
sub_regime='classification'

```

```

loader = DataSetLoader(regime=regime)

dataset = loader.load_supervised_dataset(
    filename=data_pth,
    metadata_filename=metadata_pth,
    file_type='csv')

sensitive_col_names = dataset.meta.sensitive_col_names

# use logistic regression model
model = LogisticRegressionModel()

# set the primary objective to be log loss
primary_objective = objectives.binary_logistic_loss

```

Next, let's create and save the specification files for each of the four levels of ϵ .

```

from seldonian.spec import createSupervisedSpec

# define behavioral constraints (epsilon = 0.2)
epsilon = 0.2
constraint_name = "equalized_odds"
if constraint_name == "equalized_odds":
    constraint_strs = [f'abs((FNR | [Black]) - (FNR | [White])) + abs((FPR | [Black]) - (FPR | [White])) < {epsilon}']
    deltas = [0.05]

# create spec file
save_dir = "/home/dasienga24/Statistics-Senior-Honors-Thesis/Python/COMPAS Simulation/Single Run/equalized_odds_epsilon_0.2"

os.makedirs(save_dir, exist_ok=True) #create folder

createSupervisedSpec(
    dataset=dataset,
    metadata_pth=metadata_pth,
    constraint_strs=constraint_strs,
    deltas=deltas,
    save_dir=save_dir,
    save=True,
    verbose=False)

```

```
## <seldonian.spec.SupervisedSpec object at 0x7ffe8e556650>
```

```

#-----#

# define behavioral constraints (epsilon = 0.1)
epsilon = 0.1
constraint_name = "equalized_odds"
if constraint_name == "equalized_odds":
    constraint_strs = [f'abs((FNR | [Black]) - (FNR | [White])) + abs((FPR | [Black]) - (FPR | [White])) < {epsilon}']
    deltas = [0.05]

# create spec file

```

```

save_dir = "/home/dasienga24/Statistics-Senior-Honors-Thesis/Python/COMPAS Simulation/Single Run/equali
os.makedirs(save_dir, exist_ok=True) #create folder

createSupervisedSpec(
    dataset=dataset,
    metadata_pth=metadata_pth,
    constraint_strs=constraint_strs,
    deltas=deltas,
    save_dir=save_dir,
    save=True,
    verbose=False)

```

```
## <seldonian.spec.SupervisedSpec object at 0x7ffe8de32450>
```

```

#-----#

# define behavioral constraints (epsilon = 0.05)
epsilon = 0.05
constraint_name = "equalized_odds"
if constraint_name == "equalized_odds":
    constraint_strs = [f'abs((FNR | [Black]) - (FNR | [White])) + abs((FPR | [Black]) - (FPR | [White])) < epsilon']
    deltas = [0.05]

# create spec file
save_dir = "/home/dasienga24/Statistics-Senior-Honors-Thesis/Python/COMPAS Simulation/Single Run/equali
os.makedirs(save_dir, exist_ok=True) #create folder

createSupervisedSpec(
    dataset=dataset,
    metadata_pth=metadata_pth,
    constraint_strs=constraint_strs,
    deltas=deltas,
    save_dir=save_dir,
    save=True,
    verbose=False)

```

```
## <seldonian.spec.SupervisedSpec object at 0x7ffe8d7f3a10>
```

```

#-----#

# define behavioral constraints (epsilon = 0.01)
epsilon = 0.01
constraint_name = "equalized_odds"
if constraint_name == "equalized_odds":
    constraint_strs = [f'abs((FNR | [Black]) - (FNR | [White])) + abs((FPR | [Black]) - (FPR | [White])) < epsilon']
    deltas = [0.05]

# create spec file
save_dir = "/home/dasienga24/Statistics-Senior-Honors-Thesis/Python/COMPAS Simulation/Single Run/equali
os.makedirs(save_dir, exist_ok=True) #create folder

```

```
createSupervisedSpec(
    dataset=dataset,
    metadata_pth=metadata_pth,
    constraint_strs=constraint_strs,
    deltas=deltas,
    save_dir=save_dir,
    save=True,
    verbose=False)
```

```
## <seldonian.spec.SupervisedSpec object at 0x7ffe8d8bbb90>
```

Finally, let's run the Seldonian engine for each of the four specification files.

```
from seldonian.seldonian_algorithm import SeldonianAlgorithm
from seldonian.utils.io_utils import load_pickle

# load the spec file (epsilon = 0.2)
specfile = '/home/dasienga24/Statistics-Senior-Honors-Thesis/Python/COMPAS Simulation/Single Run/equali
spec = load_pickle(specfile)
SA_02 = SeldonianAlgorithm(spec)

#-----#

# load the spec file (epsilon = 0.1)
specfile = '/home/dasienga24/Statistics-Senior-Honors-Thesis/Python/COMPAS Simulation/Single Run/equali
spec = load_pickle(specfile)
SA_01 = SeldonianAlgorithm(spec)

#-----#

# load the spec file (epsilon = 0.05)
specfile = '/home/dasienga24/Statistics-Senior-Honors-Thesis/Python/COMPAS Simulation/Single Run/equali
spec = load_pickle(specfile)
SA_005 = SeldonianAlgorithm(spec)

#-----#

# load the spec file (epsilon = 0.01)
specfile = '/home/dasienga24/Statistics-Senior-Honors-Thesis/Python/COMPAS Simulation/Single Run/equali
spec = load_pickle(specfile)
SA_001 = SeldonianAlgorithm(spec)
```

Convergence

epsilon = 0.2

```
passed_safety_02, solution_02 = SA_02.run(write_cs_logfile=True)
passed_safety_02
```

```
## True
```

epsilon = 0.1

```
passed_safety_01, solution_01 = SA_01.run(write_cs_logfile=True)
passed_safety_01
```

True

epsilon = 0.05

```
passed_safety_005, solution_005 = SA_005.run(write_cs_logfile=True)
passed_safety_005
```

True

epsilon = 0.01

```
passed_safety_001, solution_001 = SA_001.run(write_cs_logfile=True)
passed_safety_001
```

True

Accuracy

Discrimination